

1 Short Questions

Part (a) (5 pts) Give an asymptotically tight solution to the following recurrence. No justification required.

$$T(n) = T(n/4) + T(3n/4) + n^3 \quad \text{for } n \geq 4 \text{ and } T(n) = 1 \text{ for } n = 1, 2, 3.$$

Solution: Building out the recursion tree we find the work at level k is $(\frac{7}{16})^k n^3$. This results in a decreasing geometric series, which converges to $\Theta(n^3)$ ■

Part (b) (5 pts) Let $G = (V, E)$ be a *directed* graph. Given an edge $e = (u, v)$ describe a linear time algorithm to find the shortest cycle in G that contains e or report that there is no cycle containing e . The length of the cycle is simply the number of edges in it.


Solution: The shortest directed cycle containing e consists of e along with the shortest path $v \rightsquigarrow u$. Thus to find this shortest cycle (or report that none exists), we simply run *breadth-first search* from v to find the shortest path $v \rightsquigarrow u$ in $O(V + E)$ time. ■

Rubric: 10 points.

- 5 points for part (a): all or nothing.
- 5 points for part (b):
 - 4 points for the algorithm. 2 points for a superlinear algorithm.
 - 1 point for time analysis
 - -1 for a minor error, -2 for a major error.


2 Directed Graphs

Let $G = (V, E)$ be a directed graph. Each vertex $v \in V$ has a weight $w(v)$ associated with it. Given a vertex $s \in V$ let $\alpha(s) = \min\{w(v) \mid s \text{ can reach } v \text{ in } G\}$ be the minimum weight among the weights of all nodes that s can reach in G . In the figure below $\alpha(a) = 3$ and $\alpha(g) = 6$.



- List the strongly connected components in the example graph and draw the meta-graph of G .

Solution: The strongly connected components are $\{a, b, f\}, \{c\}, \{d, e\}, \{g\}, \{h\}, \{i\}$. The meta-graph is shown below:



- Suppose G is DAG. Describe a linear-time algorithm that computes $\alpha(s)$ for every $s \in V$.

Solution: Let v_1, \dots, v_n be a topological ordering of the vertices of G . For $1 \leq i \leq n$, $\alpha(v_i)$ satisfies the recurrence

$$\alpha(v_i) = \min \left\{ w(v_i), \min_{v_i \rightarrow v_j \in E} \{ \alpha(v_j) \} \right\}.$$

Note that when v_i is a sink, the second term is a min over an empty set, which we interpret to be ∞ , so no explicit base case is necessary. Since $\alpha(v_i)$ only depends on $\alpha(v_j)$ where $i < j$ (due to the topological ordering), we can memoize this into a one-dimensional array, in order from n down to 1 (i.e., in reverse topological order). We return the entire memoization structure, since we are asked to compute $\alpha(s)$ for all $s \in V$. The running time is $O(V + E)$.

COMPUTEALPHA(G):

```

 $v_1, \dots, v_n \leftarrow$  a topological ordering of the vertices of  $G$ 
for  $i$  from  $n$  down to 1
     $\alpha[v_i] \leftarrow w(v_i)$ 
    for each edge  $v_i \rightarrow v_j$ 
         $\alpha[v_i] \leftarrow \min \{ \alpha[v_i], \alpha[v_j] \}$ 
return  $\alpha$ 
```

- Extend the algorithm in the previous part to the case of a general directed graph. If you cannot figure out the previous part, you can use it as a black box in this part.

Solution: Note that for a strongly-connected component S , $\alpha(s)$ is the same for all $s \in S$, since the set $\text{rch}(G, s) = \{v \in V \mid s \text{ can reach } v\}$ is the same for all $s \in S$. Moreover, for any other vertex u , if u can reach any vertex of S then also u can reach the minimum-weight vertex of S .

Thus we take the strongly-connected component metagraph G^{SCC} , and for each metavertex \bar{v} , we set $w(\bar{v}) = \min_{u \in \text{comp}(\bar{v})} w(u)$, where $\text{comp}(\bar{v})$ is the strongly-connected component associated with \bar{v} . We then run the algorithm for the previous part on G^{SCC} with these vertex weights, and then for each vertex u in G , we will set $\alpha(u)$ to be the α -value found for the metavertex corresponding to the strongly-connected component containing u .

Forming the metagraph G^{SCC} and computing its weights takes $O(V + E)$ time, as does running the linear-time algorithm for the previous part, for a total of $O(V + E)$ time. ■

Rubric: 10 points.

- 2 points for the first part: 1 point for the vertices (i.e. the list of strongly connected components), 1 point for the edges.
- 5 points for the second part
 - 2 points max for a super-linear algorithm.
 - Scaled Dynamic programming rubric if appropriate
 - Otherwise, 4 points for the algorithm, 1 point for time analysis. −1 for minor errors, −2 for major errors.
- 3 points for the last part: 2 points for the algorithm, 1 point for time analysis. −1 for minor errors.

3 Many points further away

Let $P = \{p_1, p_2, \dots, p_n\}$ be n points in the 2-dimensional plane where each point p_i is specified as a tuple (x_i, y_i) : x_i is the x -coordinate value and y_i is the y -coordinate value for p_i . Given two points $p = (a, b)$ and $q = (c, d)$ in the plane, the L_2 distance (also called the Euclidean distance) between p and q is defined as $\sqrt{|a - c|^2 + |b - d|^2}$ (note that distance is always non-negative). Given the set P of n points and a point $q = (a, b)$, describe an $O(n)$ time algorithm to find a point $p_i \in P$ such that there are at least $n/5$ points in P that are at least as far away from q as p_i . An $O(n \log n)$ time algorithm will get you half the points (you may want to think about the slower algorithm first).

Solution (If it looks like Selection, and quacks like Selection...): We need to return a point in P whose distance from q is at most the $n/5$ -th largest (i.e., $4n/5$ -th smallest) distance from q . We do this by computing all of the distances from q to points in P , letting r be the distance of rank $4n/5$, and returning some point whose distance from q is smaller than r . Note that since distances are non-negative, it is equivalent to work with *squared* distances to avoid needing to take square roots.

```

THRESHOLDPOINT( $P[1..n]$ ,  $q = (a, b)$ ):
    allocate array  $A$  of length  $n$ 
    for  $i \leftarrow 1$  to  $n$ 
         $A[i] \leftarrow (x_i - a)^2 + (y_i - b)^2$ 
     $r \leftarrow \text{SELECT}(A[1..n], 4n/5)$ 
    for  $i \leftarrow 1$  to  $n$ 
        if  $A[i] \leq r$ 
            return  $p_i$ 

```

Using a linear-time selection algorithm (as shown in lecture), the algorithm takes $O(n)$ time. ■

Solution (...there might be a cleverer solution): The closest point to q *trivially* satisfies the requirement. The following algorithm scans through the list and finds this element in $O(n)$ time:

```

CLOSESTPOINT( $P[1..n]$ ,  $q = (a, b)$ ):
     $minDist \leftarrow \infty$ 
     $closest \leftarrow \text{NULL}$ 
    for  $i \leftarrow 1$  to  $n$ 
         $dist \leftarrow (x_i - a)^2 + (y_i - b)^2$ 
        if  $dist < minDist$ 
             $minDist \leftarrow dist$ 
             $closest \leftarrow p_i$ 
    return  $closest$ 

```

Rubric: 10 points.

- 8 points for the algorithm:
 - 5 points max for an $O(n \log n)$ algorithm.
 - 3 points max for a correct $\omega(n \log n)$ algorithm.
 - -1 for minor errors, -2 for major errors.
- 2 points for the time analysis.

4 Division

Let x and y be two positive integers with at most n digits each. Let $x \div y$ denote the maximum integer k such that $yk \leq x$. For instance $341873418723478137 \div 234334234324747 = 1458$. Describe an algorithm to compute $x \div y$ in time polynomial in n and justify its running time. You can assume that adding and subtracting (and hence comparison) of n digit numbers takes $O(n)$ time and that multiplication of two n digit numbers takes $M(n)$ time where $M(n)$ is at most $O(n^2)$. Your running time can be expressed in terms of $M(n)$. For this problem you should assume that single digit operations take $O(1)$ time but not general arithmetic operations; if you use such operations the time should be accounted for. You can assume that x and y are in binary if it helps you.

Solution: Note that $1 \leq x \div y \leq x$, so we can binary search over this space for the correct number. To divide by two, we use a bit-shift operation *shift*, that can be implemented to run in $O(n)$ time by simply copying bits over one by one. (We thus assume that x and y are given in binary.)

```

DIVIDEVIA DIVIDEANDCONQUER( $x, y$ ):
   $\ell \leftarrow 1$ 
   $r \leftarrow x$ 
  while  $\ell < r$ 
     $mid \leftarrow \text{shift}(r + \ell) \ll \lfloor \frac{r+\ell}{2} \rfloor$ 
    if  $mid \cdot y > x$ 
       $r \leftarrow mid - 1$ 
    else  $\langle mid \cdot y \leq x \rangle$ 
       $\ell \leftarrow mid$ 
  return  $\ell$ 

```

In each iteration of the while loop, we perform at most two additions ($O(n)$ each), one shift ($O(n)$), one multiplication ($M(n)$), and one comparison ($O(n)$), for a total of $O(n + M(n))$ time per iteration. The number of iterations is $O(\log_2(x)) = O(n)$, for a total running time of $O(n^2 + nM(n))$. ■

Rubric: 10 points.

- 8 points for the algorithm:
 - 6 points max for $\omega(n^2 + nM(n))$ time algorithm.
 - 2 points max for a algorithm that is exponential in n (e.g., “linear” scan)
 - −1 for a minor errors, −2 for major errors. A small number of off-by-ones collectively count as one minor error.
- 2 points for the time analysis.

5 Fire Stations

A long straight country road can be modeled as a line starting at 0. The road has n houses at locations $x_1 < x_2 < \dots < x_n$ on the line. The city wants to build fire stations along the road such that every house is within distance D from some fire station. Fire stations cannot be built at arbitrary locations. The city has figured m potential locations on the road at $y_1 < y_2 < \dots < y_m$. For simplicity assume that all the x and y values are distinct. The cost of building a fire station at location y_j is c_j . Describe an efficient algorithm that minimizes the total cost of building the fire stations with the constraint that each house is within a distance D of some fire station. Note that it is relatively easy to check whether a feasible solution exists.

Solution:

- We will add sentinel fire stations at locations $y_0 < x_1$ and $y_{m+1} = \infty$, each of cost ∞ , and a sentinel house at location $x_{n+1} = \infty$. Let $\text{MinFire}(i, j)$ denote the minimum cost to cover the houses at locations x_j, \dots, x_n using potential fire stations at locations y_i, \dots, y_m .
- We want to return $\text{MinFire}(0, 1)$.
- For $0 \leq i \leq m$, let $\text{Right}(i) = \min \{j \mid x_j > y_i + D\}$ be the index of the first house to the right of y_i that is not covered by y_i .

Then MinFire satisfies the recurrence

$$\text{MinFire}(i, j) = \begin{cases} 0 & \text{if } j > n \\ \infty & \text{if } x_j < y_i - D \\ \min \begin{cases} c_i + \text{MinFire}(i+1, \text{Right}(i)) \\ \text{MinFire}(i+1, j) \end{cases} & \text{otherwise} \end{cases}$$

(These cases correspond to: not needing to cover any more houses; the j -th house cannot be covered with the remaining stations; and choosing whether or not to open the station at y_i .)

- $\text{MinFire}(i, j)$ depends on entries of the form $(i+1, k)$. We memoize into a two-dimensional array, for i in descending order in the outer loop and for j in any order in the inner loop.
- $\text{Right}(i)$ can be precomputed for $0 \leq i \leq m$ as follows:

```

KNOWYOURRIGHTS( $x, y$ ):
   $j \leftarrow 1$ 
  for  $i \leftarrow 0$  to  $m$ 
    while  $x_j < y_i + D$ 
       $j \leftarrow j + 1$ 
     $\text{Right}[i] \leftarrow j$ 
  return  $\text{Right}$ 


```

The running time of the precomputation is $O(m + n)$ since within the nested loops, j cannot be incremented more than n times.

With this precomputation filling each entry of the memoization structure takes $O(1)$ time, for a total of $O(mn)$ time.

For completeness, we include the iterative pseudocode on the next page.

```
THISBURNINGDESIRE( $x, y, c$ ):  
   $Right \leftarrow \text{KNOWYOURRIGHTS}(x, y)$   
  for  $i \leftarrow m + 1$  down to 0  
    for  $j \leftarrow n + 1$  down to 1  
      if  $j > n$   
         $MinFire[i, j] \leftarrow 0$   
      else if  $x_j < y_i - D$   
         $MinFire[i, j] \leftarrow \infty$   
      else  
         $MinFire[i, j] \leftarrow \min \{c_i + MinFire[i + 1, Right[i]], MinFire[i + 1, j]\}$   
  return  $MinFire[0, 1]$ 
```



Rubric: 10 points. Standard Dynamic programming rubric. No penalty for slower polynomial time algorithms.

6 Shortest Paths

Let $G = (V, E)$ be a directed graph with non-negative edge lengths; $\ell(e)$ denotes the length of edge e . Suppose you have computed the shortest path distance from s to t and are not too happy about it. You have the ability to add one edge to the graph G to reduce the shortest path distance but you have to choose this edge from a given list of edges $E' = \{e_1 = (u_1, v_1), e_2 = (u_2, v_2), \dots, e_k = (u_k, v_k)\}$ where each of these edges also has its length $\ell(e_i)$ specified to you. Design an algorithm that finds the best edge to add so that the resulting graph has the smallest shortest path distance from s to t . Ideally your algorithm's running time should be $O(k)$ plus the asymptotic time to run Dijkstra's algorithm. Slower algorithms get fewer points but incorrect algorithms get few points if at all.

Solution (Graph Modeling): We will create two copies of the graph, one representing not having taken one of the edges in E' , and one representing having taken of them.

- $V^* := V \times \{\text{No}, \text{Yes}\}$
- $E^* := \{(u, a) \rightarrow (v, a) \mid u \rightarrow v \in E, a \in \{\text{No}, \text{Yes}\}\} \cup \{(u, \text{No}) \rightarrow (v, \text{Yes}) \mid u \rightarrow v \in E'\}$. For all edges, $\ell^*((u, a) \rightarrow (v, b)) = \ell(u \rightarrow v)$.
- We need to find the shortest path from (s, No) to (t, Yes) , and extract the edge in E^* corresponding to some $e_i \in E'$ that this path passes through.
- Since all edge lengths are non-negative, we will run Dijkstra's algorithm from (s, No) .
- Building $G^* = (V^*, E^*)$ by brute force takes $O(V^* + E^*) = O(V + E + k)$ time, and running Dijkstra's algorithm on G^* takes $O(E^* + V^* \log V^*) = O(k + E + V \log V)$ time. ■

Solution (Combining Shortest Paths): For each $e_i = u_i \rightarrow v_i \in E'$, the length of the shortest path going through e_i has length $\text{dist}(s, u_i) + \ell(e_i) + \text{dist}(v_i, t)$. Thus if we compute $\text{dist}(s, v)$ and $\text{dist}(v, t)$ for all $v \in V$, we can then iterate over E' to find the edge that minimizes the relevant length.

WRINKLEINTIME(G, E'):

compute $\text{dist}(s, u)$ for all u via Dijkstra's algorithm
 compute $\text{dist}(v, t)$ for all v via Dijkstra's algorithm on G^{rev}
 return $\arg \min_{e_i \in E'} \{\text{dist}(s, u_i) + \ell(e_i) + \text{dist}(v_i, t)\}$

Finding the correct edge takes $O(k)$ time, so overall the algorithm takes $O(k + E + V \log V)$ time. ■

Rubric: 10 points.

- No penalty for citing Dijkstra's algorithm as $O(E \log V)$ (and thus getting $O(k \log V)$ plus Dijkstra's).
- 8 points max for an $O(kE \log V)$ time algorithm (e.g., overly large graph construction if using graph reduction, or making $O(k)$ calls to Dijkstra's algorithm if combining multiple shortest path computations).
- 6 points max for a polynomial time algorithm slower than $O(kE \log V)$.
- For a graph reduction solution, use the Standard graph reduction rubric.
- For a combining multiple shortest path computations solution:
 - 4 points for finding the relevant shortest paths computations
 - 4 points for combining them in the correct ways
 - 2 point for time analysis in terms of the input parameters
 - -1 for each minor error, -2 for each major error